



Writing the Standards Down

Lab 2.a — the governance documents an AI agent actually uses



Two hours, mostly hands-on

6 min	Why	What shared standards buy you
18 min	The five artifacts	What each is — and how to tell them apart
10 min	How they get authored	When one is required, who reviews, drafting with Claude
26 min	Demo an ADR — then write yours	The repository boundary, documented twice
44 min	Derive + write an NFR	Judge three fixes from last week — your verdict becomes the team standard
16 min	Wire it in + review & retro	Trigger table, binding line, smoke test, reflect



Sixteen developers, sixteen private setups

The assessment called this a Cultural Gap: engineers ahead of shared norms.

WHAT WORKS TODAY

- Individual practice is genuinely strong
- Claude is productive on this repo
- The habits from 1.b landed

WHAT DOES NOT

- None of it is shared
- Every dev re-teaches Claude the same things
- No way to tell if a rule is followed

WHAT TODAY CHANGES

- Habits become written rules
- Rules get homes: documents
- Written + agreed = a standard

And the evidence says authorship matters: an ETH Zurich study (Feb 2026) found LLM-generated context files LOWERED task success ~3% while raising cost 20%+; human-written improved it ~4%. /init-and-walk-away is worse than nothing.



Every rule has a reach — that reach is its layer

A rule is the agreement; a document is its home — one document can hold many rules. A rule's LAYER is its reach: who must obey it. That decides which document it belongs in, so the right people (and agents) inherit it.

Layer	Who must obey it	A real rule at this layer	So the document lives...
Collective	Everyone, in every repo	"Never commit secrets"	In an org-wide home — which does not exist yet
Project	Everyone in ONE repo	"Responses only via the respond.ts envelope"	In this repo: CLAUDE.md, docs/adr, docs/nfr
Feature / path	Only work touching certain files	"Server tests MUST use makeTestDb()"	In .claude/rules/testing.md — loads only on matching paths

Spot the mismatch: "never commit secrets" binds every repo, but it lives in THIS repo's CLAUDE.md — a Collective-reach rule parked at Project layer, because no org-wide home exists. Building that home is the sixth artifact 2.b hands your team.



Five types — and the question each answers

Two acronyms spelled out once, then never again. If you can name the question, you know which document you are writing.

Artifact	The question it answers	What it is
CLAUDE.md	What must the agent know before anything else?	Always-loaded context + pointers
ADR	WHY is it like this?	Architecture Decision Record — one decision, with its cost
NFR	What must always be TRUE — and how do we check?	Non-Functional Requirement — a standing, testable rule
Best-practice recipe	HOW do I do X here?	A worked procedure — "recipe" from here on
Rules file	What applies ONLY in these files?	Path-scoped context (.claude/rules/)

The pair people conflate — NFR vs recipe: an NFR BINDS (what must be true, checkable; a deviation is a defect). A recipe GUIDES (the recommended route there; deviate freely, so long as the NFRs still hold).



One subject, five artifacts

This repo has a real rule: only repositories/ may import db/store.ts. Here is what each artifact says about that same subject.

Artifact	What it says about the repository boundary
CLAUDE.md	The real line, verbatim: "Don't: access db.json outside server/src/repositories/". Always loaded, no rationale
ADR	Why a repository layer was chosen, what else was considered, what it costs. Written once
NFR	The rule as a standing requirement — plus how compliance is checked
Recipe	docs/best-practices/adding-a-resource.md — the steps for adding a resource within that boundary
Rules file	Guidance that loads only when you are in the files it applies to

None of them is redundant. ADR = history. NFR = law. Recipe = procedure. CLAUDE.md = index. Rules file = context that arrives just in time.



Telling them apart when it is unclear

	CLAUDE.md	Rules file	ADR	NFR	Recipe
Answers	What to know first	What applies here	Why	What must hold	How
Tense	Present, standing	Present, scoped	Past — a choice made	Present — holds now	Imperative
Lifecycle	Revised continuously	Revised in place	Superseded, never edited	Revised in place	Updated freely
Carries a check?	No	No	No	YES — that is why	No
When it loads	ALWAYS	On matching paths	On demand, via a trigger	On demand + binding line	On demand, via a trigger
Lose it and...	The agent is unoriented	Guidance stops arriving	You lose the why	You lose the rule	You lose the how

Quick test: one-time choice → ADR · always-holds + checkable → NFR · steps → recipe · every session → CLAUDE.md · some files → rules file. One home per fact: duplicates drift — hence CLAUDE.md holds a trigger table, not the docs' contents.



How big should it get?

The scaling answer, because "keep it short" stops being useful advice at some point. The trigger table is Lab 1's lazy loading, systematized: pointer = the lazy load, trigger = when it fires.

THE HONEST ANSWER

- Lab 1.a said 100–200 lines. Right — to start.
- Past that, do not grow the file. Add a trigger row.
- The table is the mechanism: always-loaded index, on-demand detail.
- Dozens of docs later, copies drift: keep ONE home and let agents query it on demand.

CLAUDE.MD, THIS REPO

Reference docs – read on demand

When you are...	Read first	
-----	-----	
Changing response	docs/adr/0001...	
shapes or codes		
Calling an	docs/nfr/0001...	
external API		

a trigger is a SITUATION,
not a topic



Start light — and know when to graduate

This repo ships the light forms. Today you write light ADRs — and an NFR that is already one rung up, because per-clause checks are the point.

	ADR	NFR
Start here	Nygard: Context · Decision · Status · Consequences	Prose rule + one "how compliance is checked" section
Graduate to	MADR: + Decision Drivers, Considered Options	Numbered MUSTs + per-clause checks + thresholds
Graduate when	Consequences keep coming out thin	The rule starts being ignored — you want a gate, not a hope

An NFR without a verification method is a wish with a document number.



When is a document required — and who owns it?

The reason ADR practices die is not the template — nobody wrote down the trigger, the reviewer, or where it lives.

Write a...	When...	Authored / reviewed by
ADR	hard to reverse, contested, or a newcomer would ask "why is it like this?"	Anyone; peer PR review
NFR	the rule spans features AND you can say how it is checked	A few author, many adopt; EM reviews
Recipe	you have answered the same question twice	The person who answered; peer review
CLAUDE.md line	the agent needs it every session — and it is one line	Champion-gated: it costs every session
Rules file	it only applies when specific files are open	Same as a recipe

Peer review via PR is the acceptance gate today. An unreviewed standard is a draft. The "how we use these" README (2.b) exists to record exactly this table for your team.



Drafting with Claude, honestly

The ETH finding again: generated context hurts, human-written helps. Draft, do not delegate.

THE DIVISION OF LABOUR

- Claude drafts from EVIDENCE in the code.
- You own the Decision and the Consequences.
- Never ask it what the team should do.
- Read the draft as someone who disagrees.

PROMPT STARTERS (2 OF 5 — CARD IN HANDOUT)

```
# gather evidence
Read @server/src/repositories/ and
@server/src/services/todos.service.ts.
What boundary do these files enforce,
and what would break if it were
violated? Don't propose changes.
```

```
# draft, but leave the decision
Using @docs/adr/template.md, draft
an ADR for that boundary. Fill
Context and Consequences from what
you just read. Leave Decision as a
question for me to answer.
```



Keeping them true

The half of governance nobody plans for: what happens when the world moves.

- **Supersede, never edit, an accepted ADR.** The old decision happened. Write a new one that replaces it.
- **Revise NFRs in place.** They describe what holds now.
- **Every rule cites the incident that produced it.** A "Why" naming a real failure survives staff turnover; "best practice" does not.
- **Review on a cadence, not on vibes.** A date on the document, and a person against it.



Demo — write an ADR

The repository boundary. Watch the whole shape — then you write the same one.



A rule with no recorded reason

Real, load-bearing, and undocumented — the ADR trigger on all three counts.

- **The rule exists.** CLAUDE.md says routes and services never import ``db/store.ts``.
- **The code obeys it.** Only ``server/src/repositories/`` touches the store.
- **Nothing says why.** A newcomer sees a constraint with no rationale — so they route around it.
- **Nothing checks it.** Which is next week's problem, and worth noticing now.



Evidence in, decision mine

Eight minutes, live. The interesting part is the Consequences section.

WHAT I AM DOING

- Prompt Claude to READ, not to conclude.
- Paste its Context. Edit it for accuracy.
- Write the Decision myself, in one sentence.
- Add a consequence that is genuinely a cost.

MATCHES DOCS/ADR/TEMPLATE.MD

```
# docs/adr/0002-repository-only-  
#   store-access.md  
  
- Status: Accepted  
- Date: <today>  
- Deciders: <the team>  
  
## Context  
<from Claude, edited>  
  
## Decision  
Only modules under  
server/src/repositories/ may import  
db/store.ts. <-- I wrote this  
  
## Consequences  
- Good: persistence is swappable  
- Bad: one more hop for simple reads
```



Now write yours — same subject, your words

The clause list later is group work. Every document today is yours, in your copy.

- **Ten minutes, in your own copy.** Same boundary, your own Context and Consequences. The motion is the point, not novelty.
- **Use the two prompts you just watched.** Evidence first, then draft-with-the-decision-left-open. Both are in the handout.
- **Write the Decision yourself.** One sentence. If it came out of Claude, rewrite it.
- **Find one honest cost.** The acceptance bar has a cost in Consequences — "none" is not an answer.



Checkpoint 1

NOBODY ADVANCES UNTIL THIS IS TRUE

You have docs/adr/0002-*.md staged in YOUR copy: metadata block (Status / Date / Deciders) complete, all three sections — Context, Decision, Consequences — filled, and at least one real cost.

Thumbs up, or say so. If your Decision line came out of Claude rather than out of you, rewrite it now — that is the one part of an ADR that cannot be delegated.



Your turn — derive a rule

Three real fixes from last week. You decide what the standard is.



Three fixes from last week

Drawn from last week's exercise — anonymised, and where a shape didn't occur in your PRs, reconstructed. All three closed their issue. All three are green.

Fix	What shipped	CI
A	The bug is gone. No test was added.	green
B	A test was added. It passes — and it passed before the fix too.	green
C	A test was added that failed first, then passed after the fix.	green

Read them cold. No commentary from the front of the room yet.



Which are acceptable? Write the rule.

Ten minutes, in groups. You are EDITING four seed clauses, not authoring from cold — disagreement is signal.

YOUR SEED CLAUSES

- 1. A fix ships with a test
- 2. The test failed before the fix
- 3. The test lives next to the module it tests
- 4. The full suite is green before the PR opens

TAG EVERY CLAUSE

- MECHANICAL — a script decides it from the diff
- FUZZY — it needs a person
- The tag is the point. Do not skip it

TWO GROUPS

- A: the test-discipline rule (these diffs)
- B: the store boundary (only repositories/ imports db/store.ts)
- Same job: clauses, tagged

• Keep, tighten, or drop each — then
Every clause gets a verdict (keep / lighten / drop) and a tag. That is the whole task.
add your own



Where we landed

Written down before anyone moves on. This is the agreed list, not my list.

MECHANICAL

- A change under server/src ships with a test
- The test file is co-located
- Full suite green

FUZZY

- The test failed first
- The test asserts behaviour, not implementation
- The assertion is meaningful

DECIDE NOW

- Accepted risk?
- Human gate?
- Or dropped?

Merging two groups' clauses? Tiebreak: the stricter clause wins — unless someone can show it false-positives on real code in this repo. "The test failed first" is the clause that matters most, and no gate can decide it. That is the finding.



Same file, three enforcement states

A rule with half a gate. Your NFR clauses will have the same three states — that is what the tag is for.

ONE CLAUDE.MD, THREE RULES

- Store boundary: written, NOTHING checks it.
- console.log in server/src: lint FAILS it.
- logger.info with a wrong first arg: passes SILENTLY.
- Which of these is a standard? Why?

SERVER/SRC — TRY IT

```
// the rule (CLAUDE.md):  
// logger.info('<file-basename>',  
//   message, meta?)  
  
console.log('created', id)  
  // -> lint error. CAUGHT.  
  
logger.info('creating a todo',  
  { id })  
  // wrong first arg -> passes.  
  // HALF a gate.
```



What a mature team wrote

A real NFR from a mature 137-document corpus — shown now that your list exists, not before.

Section	What it contains
Classification	Category, priority, and which external framework it traces to
Requirements	Numbered MUST clauses — no prose, no hedging
Thresholds	Target / Warning / Critical, so "compliant" is measurable
Verification	Per clause: the exact test or check that proves it
Plan reference	Which phase implements it, so it is not aspirational

What did you catch that they did not? Their form is heavier than you need today; the per-clause Verification is the part worth stealing.



Turn the clauses into an NFR

Twenty minutes: your own doc, in your own copy, from the agreed clause list.

THE RULES

- Copy NFR-0001's headings and tone —
- then ADD the per-clause checks it predates.
- Every clause: "Checked by: <what>".
- Unenforceable ones get the marking convention →
- Peer review is the last five minutes. Protected.

DOCS/NFR/0002-... (YOURS)

The rules

1. A change under server/src ships with a test in the same PR.
Checked by: CI diff check.
2. The test failed before the fix.
Checked by: — NOT MECHANICALLY
CHECKABLE — decision: human gate
(reviewer asks) — <name>

clause + check + mark + name.
an unmarked unenforceable clause
is the dangerous kind.



Checkpoint 2

NOBODY ADVANCES UNTIL THIS IS TRUE

Your NFR names how each clause is checked — and the ones that cannot be are marked, with a decision and a name.

Artifact check, not a state check: I am reading your document, not your test output. A clause with no check and no mark is the one thing that fails here.



A document nothing points at is a document nobody reads

Claude included. This step is why the last two hours matter.

TWO KINDS OF WIRING

- A TRIGGER ROW per doc — fetches the full doc when the situation applies.
- A BINDING LINE for the NFR — one line, always in force, never fetched.
- Same pair as 1.a: trigger row = the lazy pointer; binding line = a one-line EAGER load.
- The doc is the law; the binding line is the sign on the wall.
- Then smoke-test in a FRESH session.

CLAUDE.MD

When you are...	Read first	
-----	-----	
Changing the store	docs/adr/0002...	
boundary		
Fixing a bug or	docs/nfr/0002...	
adding behaviour		

Do / Don't

- Do: every change under server/src ships with a test in the same PR (NFR-0002) <-- the binding line

smoke test, fresh session:

"what is our rule about tests?"



Compare, reflect, capture

Five minutes that make the difference between an exercise and a practice.

- **Two volunteers, two NFRs.** Same clause list, different documents — what did each choose to check, and what did they mark?
- **One retro question.** What was hard, what was easy? Thirty seconds each, around the room.
- **Reflection note — now, not later.** 3–5 sentences in your PR description: what you learned, what surprised you, what you'd apply tomorrow. It is an acceptance criterion.
- **Unfinished work finishes async.** Still required, still peer-reviewed — the list does not shrink because the clock ran out.



Ship checklist, and what is next

Questions & support: #extra-duty-solutions-ai-training — your instructor and the Applied AI Engineer are there.

- **Ship checklist.** ADR (cost included) · NFR (checks + marks) · wired into CLAUDE.md (trigger rows + binding line) · smoke-tested · PR-reviewed, CI green · reflection note in the PR.
- **Stretch, self-paced.** Store-boundary clauses · the half-gate ESLint clause · a rules-file conversion · a head start on the "how we use these" README.
- **Everything is in the handout.** All five prompt starters, the stuck-points table, the skeletons.
- **Next week: teeth.** Your clause becomes a hook and a CI gate — and the marked clauses get their reckoning.